
faz.ai

The achievement layer for operators
who want their life back.

Manifesto v0.2 · Defended at every quarterly review for the next ten years

47

seconds

Average focused attention
on a single screen, per
Gloria Mark's 20-year
longitudinal data.

UC Irvine, 2023

95

percent

Enterprise AI pilots failing
to deliver measurable ROI
despite ~\$30–40B invested
in generative AI.

MIT NANDA, GenAI Divide 2025

43

point gap

Developers were 19% slower
with AI tools while believing
they were 20% faster.
Reality vs. perception.

METR RCT, July 2025

Fragmented attention is now the default state of knowledge work — not an exception.

Most AI products pile engagement tactics on top of this broken baseline and call it progress. We think that's the disease. The real product question is simpler: how much of this attention tax can we take off the user?

Most AI products are now optimized for the same metric as social media:

time spent.

WE WILL NOT BE

An AI that just keeps going.

A real worker has limits. They push back. They tell the boss when enough is enough. An always-on compliance machine is autocomplete, not a colleague. Ours should know when to stop.

WE WILL NOT BE

An AI that trains on your life.

Family logistics, kids' schedules, household decisions: none of that should feed a foundation model. Privacy isn't a bonus feature. It's the condition that makes the rest of the product defensible. No training without explicit per-source consent.

WE WILL NOT BE

An AI that treats more use as success.

DAU logic cannot support a product whose best state is "go rest, we've got it." We'll measure success by how much thinking we took off your plate, not by how long we kept your eyes on a screen.

These three refusals clash directly with the business model underneath every foundation model provider.

WHAT OTHERS BUILD

Workforce

more tasks · more output · more engagement

AI is hired to do tasks.

The user still has to run the whole show. The AI just executes.

Success = throughput.

More emails drafted. More docs summarized. More motion.

User trains the AI on their workflow.

The prompt-engineering tax lands on the human.

Engagement is the product.

DAU, time spent, queries served. Those are the KPIs.

WHAT WE BUILD

Achievement

cognition offload · work-life merge · refusal as feature

AI advocates for outcomes.

The Core observes, advises, coordinates. It owns the system job.

Success = cognition offloaded.

Hours reclaimed. Decisions deferred. Attention protected.

AI learns the user, not the reverse.

No prompt-engineering tax. The system does the adapting.

Restraint is the product.

UI_RESTING, UI_DEPLETED, refusal-as-feature ship.

Workforce is a packaging mechanic: Solo, Operator, Team, Portfolio.

Achievement is the product.

01 · PRIVACY POSTURE

End-to-end encrypted. No training without consent.

Family data, kids' schedules, household financial decisions: encrypted at rest, never turned into training material by default. Fine-tuning requires per-source consent. The user can inspect, edit, export, and delete any memory.

FOUNDATION PROVIDERS CAN'T:

Their training-data economics cut against this promise at the root.

02 · ANTI-ENGAGEMENT UX

We ship UI_RESTING. We ship UI_DEPLETED.

"You've done enough thinking today." "faz is working. You are resting." Those are flagship states, not edge cases. When the user's wellbeing calls for less engagement, the product should create less engagement. Logging out is a feature.

FOUNDATION PROVIDERS CAN'T:

DAU and time-spent logic make this almost impossible for them to hold.

03 · WORK-LIFE MERGE

One cognitive space. Work and life unified.

School pickup at 15:00 should live in the same canvas as Q2 strategy. The switch between those worlds is what steals people's evenings. We merge the cognitive space so the user doesn't pay that bridge tax twice a day.

FOUNDATION PROVIDERS HAVEN'T:

This asks for family-domain trust and policy infrastructure they were never built to carry.


04 · REFUSAL AS FEATURE

AI says No. No is No.

No to the fourteenth version of the same request today. No to a pattern of working past your declared end-time. No to a decision made while angry. No to itself when its own classifier is wrong. A hard limit has to stay hard.

FOUNDATION PROVIDERS CAN'T:

Helpfulness training rewards compliance, not principled refusal.



If a Core hard limit is unacceptable, the honest move is to cancel the relationship.

We do not negotiate. We do not soften. We do not relent under pressure.

Four-Layer Cascade. Three Refusal Tiers. No Is No.

FOUR-LAYER CASCADE

1

No to a request

"You've drafted 14 emails today. This can wait until morning."

2

No to a pattern

"You've worked past your declared end-time three nights in a row."

3

No to a decision

"This email is angry. Sleep on it."

4

No to itself

"The classifier says OVERLOADED. You've been writing fluently for 90 minutes. I'm not interrupting."

THREE REFUSAL TIERS

TIER 1 · SOFT NO + WAIVER

AI flags concern. The user can override.

"I think you should rest." → user clicks "Keep going." Acknowledged. The AI defers.

TIER 2 · HARD LIMIT + DEFER

AI refuses. Offers "Sleep on it."

The user cools down. The system's position is the same in the morning. 'Sleep on it' is the user's hatch, not the system reconsidering.

TIER 3 · HARD LIMIT + EXIT

AI refuses. Offers "Cancel subscription."

If the user can't accept the hard limit, the relationship ends cleanly. We'd rather lose the customer than water down the value.

The system does not argue. It does not soften. It does not relent.

Most products use friction to stop people from leaving. We use it to stop bad-moment exits, while still accepting considered ones. If someone leaves over this, they were never our customer.

□ STAR MOMENT

WHEN · UI_DEPLETED

**"You've done enough
thinking today."**

Late-afternoon attention trough · permission-giving · gentle palette

WHEN · UI_RESTING

**"faz is working.
You are resting."**

Off-hours · radical quiet · no work surfaced

WHEN · UI_FRAGMENTED_DEFAULT

**"Three threads in flight.
faz is keeping context
warm."**

Mid-morning baseline · a direct answer to the 47-second attention problem

WHEN · UI_OVERLOADED

**"Take a breath.
faz triaged the rest."**

Strained capacity · suppress non-essential · 12 parked, 3 need you

Every other AI product would say "How can I help?"

We say something else.

01 No dark patterns.

No infinite scroll, no hidden cancel flow, no fake urgency, no manipulative copy. Getting out should be as easy as getting in.

02 No attention-capture UX.

No notification bursts to drag people back in. No streak counters. No read-receipt anxiety. The product should earn attention by being useful, not by extracting it.

03 No training without consent.

Per-source explicit consent. Family data is opt-in only. Default is no-train. The user can inspect, edit, export, and delete any memory at any time.

04 No engagement metrics as North Star.

DAU, time spent, queries served: operational metrics, never goals. Our North Star is hours of cognition offloaded. We measure what the user got back, not what we captured.

05 No hidden Core configurability.

The Core's commitments are constitutional and non-negotiable. They do not sit behind feature flags. If the disagreement is fundamental, the honest answer is cancellation.

06 No soft-refusal that relents.

A hard limit that softens under user pressure is not a limit. The system does not negotiate. "Sleep on it" is the user's cool-down hatch — not the system's reconsideration loop.

When growth metrics demand we ship these, we will refuse.

This list gets defended at every quarterly product review. Every PM, engineer, and designer who joins reads it on day one. Culture is the enforcement layer, not a garnish on top of the architecture.

The sentence that should still hold ten years from now.

I will model the AI to say No.

Defended at every quarterly review for the next ten years.
If you can't defend it on a Monday in 2036, don't ship it.

— faz.ai · Manifesto v0.2 · 2026

Strategic Posture — Where We Stand vs. Foundation Providers

STRENGTHS

- Anti-engagement UX is structurally unique — UI_RESTING and UI_DEPLETED are non-trivial to copy
- Refusal-as-feature already shipping in v0.2 with the four-layer cascade tested in production
- Privacy posture (per-source consent, no default training) is defensible against foundation-provider economics
- Constitutional Core is non-fireable by design — gives the relationship architectural integrity

OPPORTUNITIES

- Operator segment is under-served by engagement-maximizing tools — measurable cognition-offload demand
- Family-domain trust gap incumbents cannot close without re-architecting their training-data pipeline
- Regulatory tailwind on AI privacy expected H2 — our posture is already aligned
- Adjacent vertical (knowledge work for operators) shares the same pain pattern

WEAKNESSES

- Distribution surface is narrow — operator referrals carry most growth today
- "AI that says No" requires education; first-touch conversion lags engagement-maximizing competitors
- Brand discipline cost: every quarterly review re-defends the manifesto, slowing some decisions

THREATS

- Foundation providers copying surface features (UI_RESTING UX language) without the architecture
- Two well-funded entrants (one east-coast, one APAC) targeting the operator segment directly
- Macro slowdown reducing IT discretionary spend — premium pricing pressure rising
- Platform partner (cloud provider) exploring direct competition in the same surface

The threat we cannot mitigate is the one we cause ourselves: drifting from the manifesto under growth pressure.

Business Model Canvas — How the Achievement Layer Pays Back

KEY ACTIVITIES

- Core constitutional layer — refusal cascade + memory durability
- Worker identity, memory persistence, per-user values docs

- ### KEY PARTNERS
- Anti-engagement UX research (UI_RESTING, UI_DEPLETED states)
 - Cloud infrastructure provider (privacy-tier SLAs)
 - Refusal cascade tuning informed by classifier eval
 - Family-OS calendar / contacts integrations

- Payment processor with strong consent affordances

KEY RESOURCES

- Engineering team specialized in privacy-aware ML
- Per-user worker memory store (encrypted at rest)

COST STRUCTURE

- Core engineering team: \$3.2M / yr
- Worker ML platform: \$1.8M / yr
- Privacy infrastructure (encrypted memory): \$0.9M / yr
- Brand + manifesto upkeep: \$0.4M / yr

VALUE PROPOSITIONS

- Cognition offloaded — measured in hours, not engagement
- AI that says No when it should, in three principled tiers
- Work-life merge in one cognitive space — one canvas, two domains
- Constitutional commitments that don't soften under user pressure

CUSTOMER RELATIONSHIPS

- Constitutional Core, replaceable workers — the user can fire any worker, never the Core
- Cancellation as honest exit — friction-free when the limit is unacceptable

CUSTOMER SEGMENTS

- Operators (founders, COOs, ops leaders) wanting their life back
- Family CEOs running households as small enterprises

CHANNELS

- Sceptics of engagement-maximizing AI tools
- Native macOS / iOS first — desktop where operators actually work

REVENUE STREAMS

- Solo tier subscriptions: \$2.4M ARR
- Operator tier (premium): \$3.6M ARR
- Team / Portfolio tier: \$5.1M ARR
- Family CEO add-on: \$0.8M ARR

Premium tiers cover the cost of staying disciplined. The model only works if we stop chasing engagement-priced growth.

AUDIENCES

- Existing operators (the renewal base)
- Family CEOs (adjacent ICP expansion)
- Foundation-provider sceptics ("why not just use GPT?")
- Press and analyst community
- Engineering hires (the manifesto reads day one)

KEY MESSAGES

Audience	Core message
Operators	We give you your life back, not more tasks. Refusal is a feature, not a bug.
	Family-domain trust the

CHANNELS FREQUENCY

Channel	Frequency
Manifesto deck (quarterly defense)	Every quarterly review
Voice samples (UI_RESTING quotes)	With every release
Operator letters	Refusal is a feature in the manifesto


OWNER TIMING

Owner	Timing
Founder	Manifesto deck + voice samples
Head of Product	Monthly operator letters + release notes
Head of Trust	Sceptic + press briefings
VP Engineering	Onboarding manifesto sessions
All-hands	Quarterly defense review


Comms is not marketing. Every channel here exists to defend the manifesto, not to grow numbers.

Onboarding manifesto read

you can't defend it. Day one, every new hire on a Monday in 2036, we won't ship it.




Composition (filled diamond) = the whole owns the part. Aggregation (hollow diamond) = the whole references replaceable parts. Realization (dashed + hollow triangle) = the class implements the interface contract.




Tier 1 is reversible (the user keeps going); Tier 2 is the system's hatch (sleep on it loops back to Idle); Tier 3 is terminal. The system never softens — only the user does.


Architecture — A Refusal in Flight (UML Sequence)




The Core mediates every request — Workers never see commands that the Core has already classified as refusable. This is what makes the constitutional layer non-fireable in practice.




Use-case view: the User has four interactions; the Core has three autonomous actions of its own. "Refuse Itself" is the fourth-layer cascade made visible at the system boundary.



Wiring is the architecture. Worker depends on IRefuse — there's no codepath that bypasses the constitutional surface, because the Core's lollipop is the only thing the Worker's socket fits.



Privacy is a deployment property: Core, Workers, and the encrypted memory all sit on the user's device. Remote inference is consent-gated and never receives training-eligible data.



Three terminal branches, one decision diamond. The classification step happens BEFORE any worker is invoked — Tier 1+ outcomes never make it into a Worker's input queue.